



Red Hat

Start with XDP

Hangbin Liu <haliu@redhat.com>

Software Engineering

Agenda:

- What's & Why XDP
- Write a small program
- Extend the program
- Use BPF maps
- Load the BPF program with libbpf

XDP(eXpress Data Path):

- Easy to use
- Fast

| 2. Write a small program

```
$ cat xdp_drop.c
#include <linux/bpf.h>
#include <bpf/bpf_helpers.h>

SEC("xdp_drop")
int xdp_drop_prog(struct xdp_md *ctx)
{
    return XDP_DROP;
}

char _license[] SEC("license") = "GPL";
```

| 2.1: Build eBPF/XDP object

```
$ sudo dnf install clang llvm gcc kernel-headers
$ sudo dnf install libbpf-devel libxdp-devel xdp-tools bpftool
$ clang -O2 -g -Wall -target bpf -c xdp_drop.c -o xdp_drop.o
$ llvm-objdump -S -no-show-raw-insn xdp_drop.o
```

```
xdp_drop:          file format ELF64-BPF
```

Disassembly of section xdp_drop:

```
0000000000000000 xdp_drop_prog:
;          return XDP_DROP;
    0:          r0 = 1
    1:          exit
```

| 2.1: Build eBPF/XDP object

```
$ llvm-objdump -h xdp_drop.o
```

```
xdp_drop:      file format ELF64-BPF
```

Sections:

| Idx | Name | Size | VMA | Type |
|-----|-----------------|----------|------------------|------|
| 0 | | 00000000 | 0000000000000000 | |
| 1 | .strtab | 000000ad | 0000000000000000 | |
| 2 | .text | 00000000 | 0000000000000000 | TEXT |
| 3 | xdp_drop | 00000010 | 0000000000000000 | TEXT |
| 4 | license | 00000004 | 0000000000000000 | DATA |
| 5 | .debug_str | 00000125 | 0000000000000000 | |
| 6 | .debug_abbrev | 000000ba | 0000000000000000 | |
| 7 | .debug_info | 00000114 | 0000000000000000 | |
| 8 | .rel.debug_info | 000001c0 | 0000000000000000 | |
| 9 | .BTF | 000001df | 0000000000000000 | |
| 10 | .rel.BTF | 00000010 | 0000000000000000 | |
| 11 | .BTF.ext | 00000050 | 0000000000000000 | |
| 12 | .rel.BTF.ext | 00000020 | 0000000000000000 | |
| 13 | .eh_frame | 00000030 | 0000000000000000 | DATA |
| 14 | .rel.eh_frame | 00000010 | 0000000000000000 | |
| 15 | .debug_line | 00000084 | 0000000000000000 | |
| 16 | .rel.debug_line | 00000010 | 0000000000000000 | |
| 17 | .llvm_addrsig | 00000002 | 0000000000000000 | |
| 18 | .symtab | 000002d0 | 0000000000000000 | |

| 2.2: Load XDP object via ip link

```
# Warn: dont load it on default interface!
$ sudo ip link set veth1 xdpgeneric obj xdp_drop.o sec xdp_drop
$ sudo ip link show veth1
6: veth1@veth0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 xdpgeneric qdisc noqueue ...
    link/ether f6:03:7f:e7:27:bb brd ff:ff:ff:ff:ff:ff
    prog/xdp id 1 tag 57cd311f2e27366b jited

$ sudo bpftool prog show
1: xdp tag 57cd311f2e27366b gpl
    loaded_at 2021-02-04T04:12:00-0500 uid 0
    xlated 16B jited 40B memlock 4096B

$ sudo xdp-loader status
CURRENT XDP PROGRAM STATUS:
```

| Interface | Prio | Program name | Mode | ID | Tag | Chain actions |
|-----------|------|------------------|------|----|------------------|---------------|
| lo | | <no XDP program> | | | | |
| ens3 | | <no XDP program> | | | | |
| veth1 | | | skb | 1 | 57cd311f2e27366b | |

```
$ sudo ip link set veth1 xdpgeneric off
```

| 2.2: Load bpf object via xdp-loader(recommend)

```
# Warn: dont load it on default interface!
$ sudo xdp-loader load -m skb -s xdp_drop veth1 xdp_drop.o
$ sudo xdp-loader status
CURRENT XDP PROGRAM STATUS:

Interface      Prio  Program name      Mode   ID   Tag               Chain
actions
-----
lo             <no XDP program>
ens3          <no XDP program>
veth1         xdp_dispatcher    skb    15   d51e469e988d81da
=>            50   xdp_drop_prog      20   57cd311f2e27366b XDP_PASS

$ sudo ip link show veth1
4: veth1@if3: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 xdpgeneric qdisc noqueue ...
    link/ether ba:4d:98:21:3b:b3 brd ff:ff:ff:ff:ff:ff link-netns ns
    prog/xdp id 15 tag d51e469e988d81da jited

$ sudo bpftool prog show
15: xdp   name xdp_dispatcher tag d51e469e988d81da gpl
        loaded_at 2021-01-13T03:24:43-0500 uid 0
        xlated 616B jited 638B memlock 4096B map_ids 8
        btf_id 12
20: ext   name xdp_drop_prog tag 57cd311f2e27366b gpl
        loaded_at 2021-01-13T03:24:43-0500 uid 0
        xlated 16B jited 40B memlock 4096B
        btf_id 16

$ sudo xdp-loader unload -a veth1
```


| 3: Extend the program, Drop specific packets

```
/* user accessible metadata for XDP packet hook
 * new fields must be added to the end of this structure
 */
struct xdp_md {
    __u32 data;
    __u32 data_end;
    __u32 data_meta;
    /* Below access go through struct xdp_rxq_info */
    __u32 ingress_ifindex; /* rxq->dev->ifindex */
    __u32 rx_queue_index; /* rxq->queue_index */

    __u32 egress_ifindex; /* txq->dev->ifindex */
};
```

| 3: Extend the program, Drop specific packets

```
#include <linux/bpf.h>
#include <bpf/bpf_helpers.h>
#include <linux/if_ether.h>
#include <arpa/inet.h>

SEC("xdp_drop")
int xdp_drop_prog(struct xdp_md *ctx)
{
    void *data_end = (void *) (long) ctx->data_end;
    void *data = (void *) (long) ctx->data;
    struct ethhdr *eth = data;
    __u16 h_proto;

    if (data + sizeof(struct ethhdr) > data_end)
        return XDP_DROP;

    h_proto = eth->h_proto;

    if (h_proto == htons(ETH_P_IPV6))
        return XDP_DROP;

    return XDP_PASS;
}

char _license[] SEC("license") = "GPL";
```

| 4: Use BPF map: count the processed packets

```
struct {
    __uint(type, BPF_MAP_TYPE_PERCPU_ARRAY);
    __type(key, __u32);
    __type(value, long);
    __uint(max_entries, 1);
} rxcnt SEC(".maps");

enum bpf_map_type {
    BPF_MAP_TYPE_UNSPEC,
    BPF_MAP_TYPE_HASH,
    BPF_MAP_TYPE_ARRAY,
    BPF_MAP_TYPE_PROG_ARRAY,
    BPF_MAP_TYPE_PERF_EVENT_ARRAY,
    BPF_MAP_TYPE_PERCPU_HASH,
    BPF_MAP_TYPE_PERCPU_ARRAY,
    [...]
    BPF_MAP_TYPE_TASK_STORAGE,
};
```

| 4: Use BPF map: count the processed packets

```
#include <linux/bpf.h>
#include <bpf/bpf_helpers.h>
#include <linux/if_ether.h>
#include <arpa/inet.h>

struct {
    __uint(type, BPF_MAP_TYPE_PERCPU_ARRAY);
    __type(key, __u32);
    __type(value, long);
    __uint(max_entries, 1);
} rxcnt SEC(".maps");

SEC("xdp_drop_ipv6")
int xdp_drop_ipv6_prog(struct xdp_md *ctx)
{
    void *data_end = (void *) (long) ctx->data_end;
    void *data = (void *) (long) ctx->data;
    struct ethhdr *eth = data;
    __u16 h_proto;
    __u32 key = 0;
    long *value;

    if (data + sizeof(struct ethhdr) > data_end)
        return XDP_DROP;

    h_proto = eth->h_proto;

    if (h_proto == htons(ETH_P_IPV6)) {
        value = bpf_map_lookup_elem(&rxcnt, &key);
        if (value)
            *value += 1;
        return XDP_DROP;
    }

    return XDP_PASS;
}

char _license[] SEC("license") = "GPL";
```

| 4.1: Load and show map counts

```
$ sudo xdp-loader load -m skb -s xdp_drop_ipv6 veth1 xdp_drop_ipv6_count.o
```

<receive some IPv6 packets>

```
$ sudo bpftool map show
```

```
bpftool map show
```

```
13: percpu_array name rxcnt flags 0x0
```

```
    key 4B value 8B max_entries 1 memlock 4096B
```

```
    btf_id 20
```

```
19: array name xdp_disp.rodata flags 0x480
```

```
    key 4B value 84B max_entries 1 memlock 8192B
```

```
    btf_id 28 frozen
```

```
# bpftool map dump id 13
```

```
[{
  "key": 0,
  "values": [{
    "cpu": 0,
    "value": 13
  }, {
    "cpu": 1,
    "value": 7
  }, {
    "cpu": 2,
    "value": 0
  }, {
    "cpu": 3,
    "value": 0
  }
}]
]
```

| 5: Load the BPF program with libbpf/libxdp

```
#include <unistd.h>
#include <stdlib.h>
#include <string.h>
#include <linux/if_link.h>
#include <signal.h>
#include <net/if.h>
#include <assert.h>

/* In this example we use libbpf-devel and libxdp-devel */
#include <bpf/bpf.h>
#include <bpf/libbpf.h>
#include <xdp/libxdp.h>

/* We define the following variables global */
static int ifindex;
struct xdp_program *prog = NULL;

/* This function will remove XDP from the link when program exit */
static void int_exit(int sig)
{
    xdp_program__close(prog);
    exit(0);
}
```

| 5: Load the BPF program with libbpf/libxdp

```
/* This function will count the per CPU number and print out total
 * dropped packets number and PPS(packets per second).
 */
static void poll_stats(int map_fd, int interval)
{
    int ncpus = libbpf_num_possible_cpus();
    if (ncpus < 0) {
        printf("Error get possible cpus\n");
        return;
    }
    long values[ncpus], prev[ncpus], total_pkts;
    int i, key = 0;

    memset(prev, 0, sizeof(prev));

    while (1) {
        long sum = 0;

        sleep(interval);
        assert(bpf_map_lookup_elem(map_fd, &key, values) == 0);
        for (i = 0; i < ncpus; i++)
            sum += (values[i] - prev[i]);
        if (sum) {
            total_pkts += sum;
            printf("total dropped %10llu, %10llu pkt/s\n",
                   total_pkts, sum / interval);
        }
        memcpy(prev, values, sizeof(values));
    }
}
```

| 5: Load the BPF program with libbpf/libxdp

```
int main(int argc, char *argv[])
{
    int prog_fd, map_fd, ret;
    struct bpf_object *bpf_obj;

    if (argc != 2) {
        printf("Usage: %s IFNAME\n", argv[0]);
        return 1;
    }

    ifindex = if_nametoindex(argv[1]);
    if (!ifindex) {
        printf("get ifindex from interface name failed\n");
        return 1;
    }

    /* load XDP object by libxdp */
    prog = xdp_program__open_file("xdp_drop_ipv6_count.o", "xdp_drop_ipv6", NULL);
    if (!prog) {
        printf("Error, load xdp prog failed\n");
        return 1;
    }
}
```


| 5: Load the BPF program with libbpf/libxdp

```
/* attach XDP program to interface with skb mode
 * Please set ulimit if you got an -EPERM error.
 */
ret = xdp_program__attach(prog, ifindex, XDP_MODE_SKB, 0);
if (ret) {
    printf("Error, Set xdp fd on %d failed\n", ifindex);
    return ret;
}

/* Find the map fd from bpf object */
bpf_obj = xdp_program__bpf_obj(prog);
map_fd = bpf_object__find_map_fd_by_name(bpf_obj, "rxcnt");
if (map_fd < 0) {
    printf("Error, get map fd from bpf obj failed\n");
    return map_fd;
}

/* Remove attached program when program is interrupted or killed */
signal(SIGINT, int_exit);
signal(SIGTERM, int_exit);

poll_stats(map_fd, 2);

return 0;
}
```

```
# Warn: dont load it on default interface!
$ sudo ulimit -l unlimited
$ gcc xdp_drop_ipv6_count_user.c -o xdp_drop_ipv6_count -lbpf -lxdp
$ sudo ./xdp_drop_ipv6_count veth1
total dropped          2,          1 pkt/s
total dropped        129,         63 pkt/s
total dropped        311,         91 pkt/s
total dropped        492,         90 pkt/s
total dropped        674,         91 pkt/s
total dropped        856,         91 pkt/s
total dropped       1038,         91 pkt/s
^C
```

All the example code could be find in

https://github.com/liuhangbin/xdp_examples

Q

&

A

Thanks!